



Controlled Re-execution for Disjunctive Predicates

EE 382N Distributed Systems

Term Project

Fall 2015

Elie Antoun

Karim Serhan

Motivation

Mo • ti • va • tion
/ˌmōdəˈvāSH(ə)n/

The driving force by
which humans
achieve their goals

Non-deterministic
nature of distributed
computations



Present bugs are
difficult to reproduce

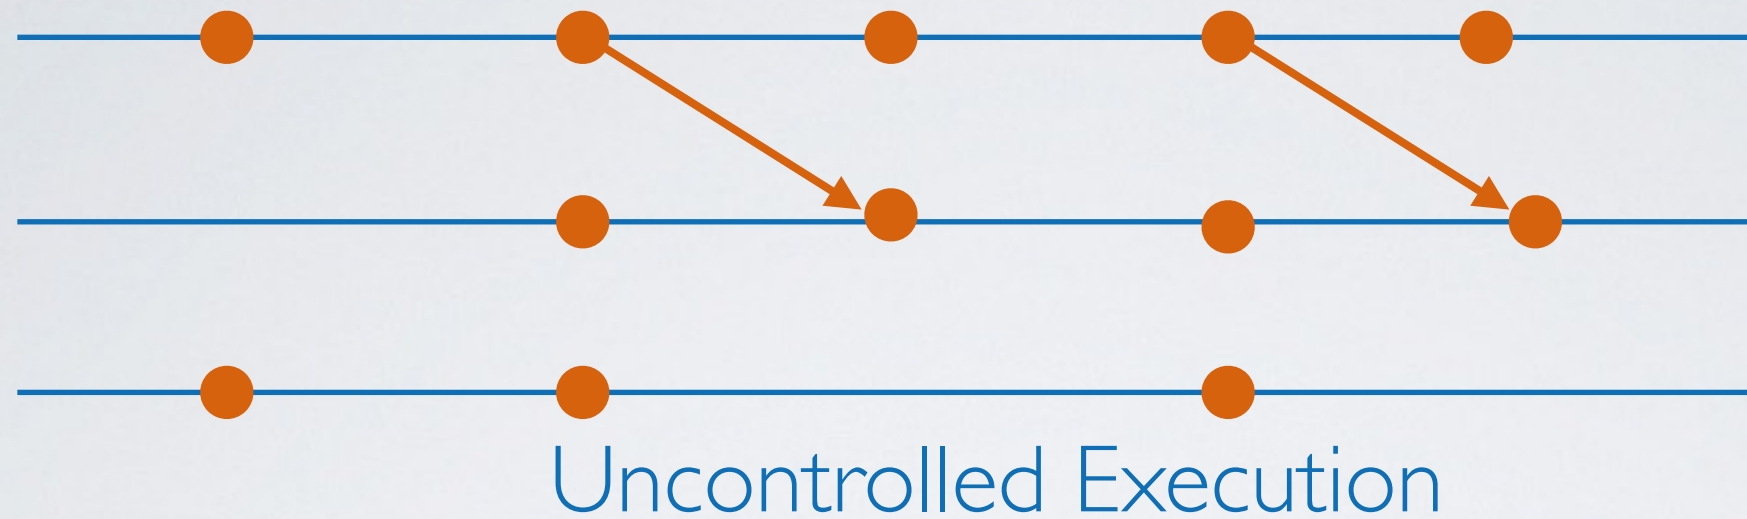


Need to control a
distributed computation


Proliferation of
distributed
computations



Control Re-execution




→
Messages between
events



→
Messages between
events
→
Synchronization
messages

Disjunctive Predicates



Global predicate B is a disjunction of local predicates on processes
 $B = L_1 \vee L_2 \vee L_3$

Problem Statement

Given a distributed computations and a predicate B , we aim to control the execution by adding synchronization messages to the computation so that B is always satisfied.

$B = \text{at least one process is alive at any time} \Rightarrow B = \text{alive}(P_1) \vee \text{alive}(P_2) \dots \vee \text{alive}(P_n)$

$B = \text{mutual exclusion for two processes} \Rightarrow B = \neg \text{CS}(P_1) \vee \neg \text{CS}(P_2)$

Admissible Sequence

For a given computation C and a predicate P , an admissible sequence of events α captures some safe execution of events S , where $S \subseteq E$

Properties of admissible sequence


1. Agreement: an admissible sequence α should be consistent with the happened before relation \rightarrow
$$\forall i, j : i < j \Rightarrow \alpha_j \nrightarrow \alpha_i$$
2. Boundary Condition: an admissible sequence must start with an initial event and end in a final event
$$(\alpha_0 \in E.\perp) \wedge (\alpha_{n-1} \in E.T)$$
4. Continuity: an admissible sequence must be continuous
$$\forall (\alpha_j, \alpha_i \in \alpha) \wedge (\alpha_j, \alpha_i \notin E.T) \wedge (\alpha_i.succ \rightarrow \alpha_j) \Rightarrow \alpha_{i+1} = \alpha_i.succ$$
5. Safety: α must be a safe execution.
$$\text{legal}(C, E, \alpha) \wedge (\exists \alpha_i \in \alpha \wedge \alpha_i \in \text{frontier}(C, E)) \Rightarrow C(B) = \text{True}$$

True Event Graph


For a given computation (E, \rightarrow) and a predicate B , a true event graph is defined as:

1. An event e belongs to the true event graph if the local predicate for the process $e.procc$ is true at the event e .
2. There exists an edge between two events e and f if $e.succ \rightarrow f$
3. An event e in the true event graph is labeled as initial if $e \in E.\perp$
4. An event e in the true event graph is labeled as “final: if $e \in E.T$

True Event Graph



Computation (E, \rightarrow)



True Event Graph

Theorems

Theorem 1

A predicate B is controllable in a computation (E, \rightarrow) iff there exists an admissible sequence of events with respect to the predicate B and the computation (E, \rightarrow)

Theorem 2


Let $G(v, E)$ be the true event graph corresponding to the computation (E, \rightarrow) and a predicate B . The shortest path in G , if it exists, corresponds to the admissible sequence of events in the computation

Theorem 3


Let $G(v, E)$ be the true event graph corresponding to the computation (E, \rightarrow) and a predicate B . B is controllable if in the computation (E, \rightarrow) then there exists a permissible path in G .

Solution


Computation Trace



True Event Graph

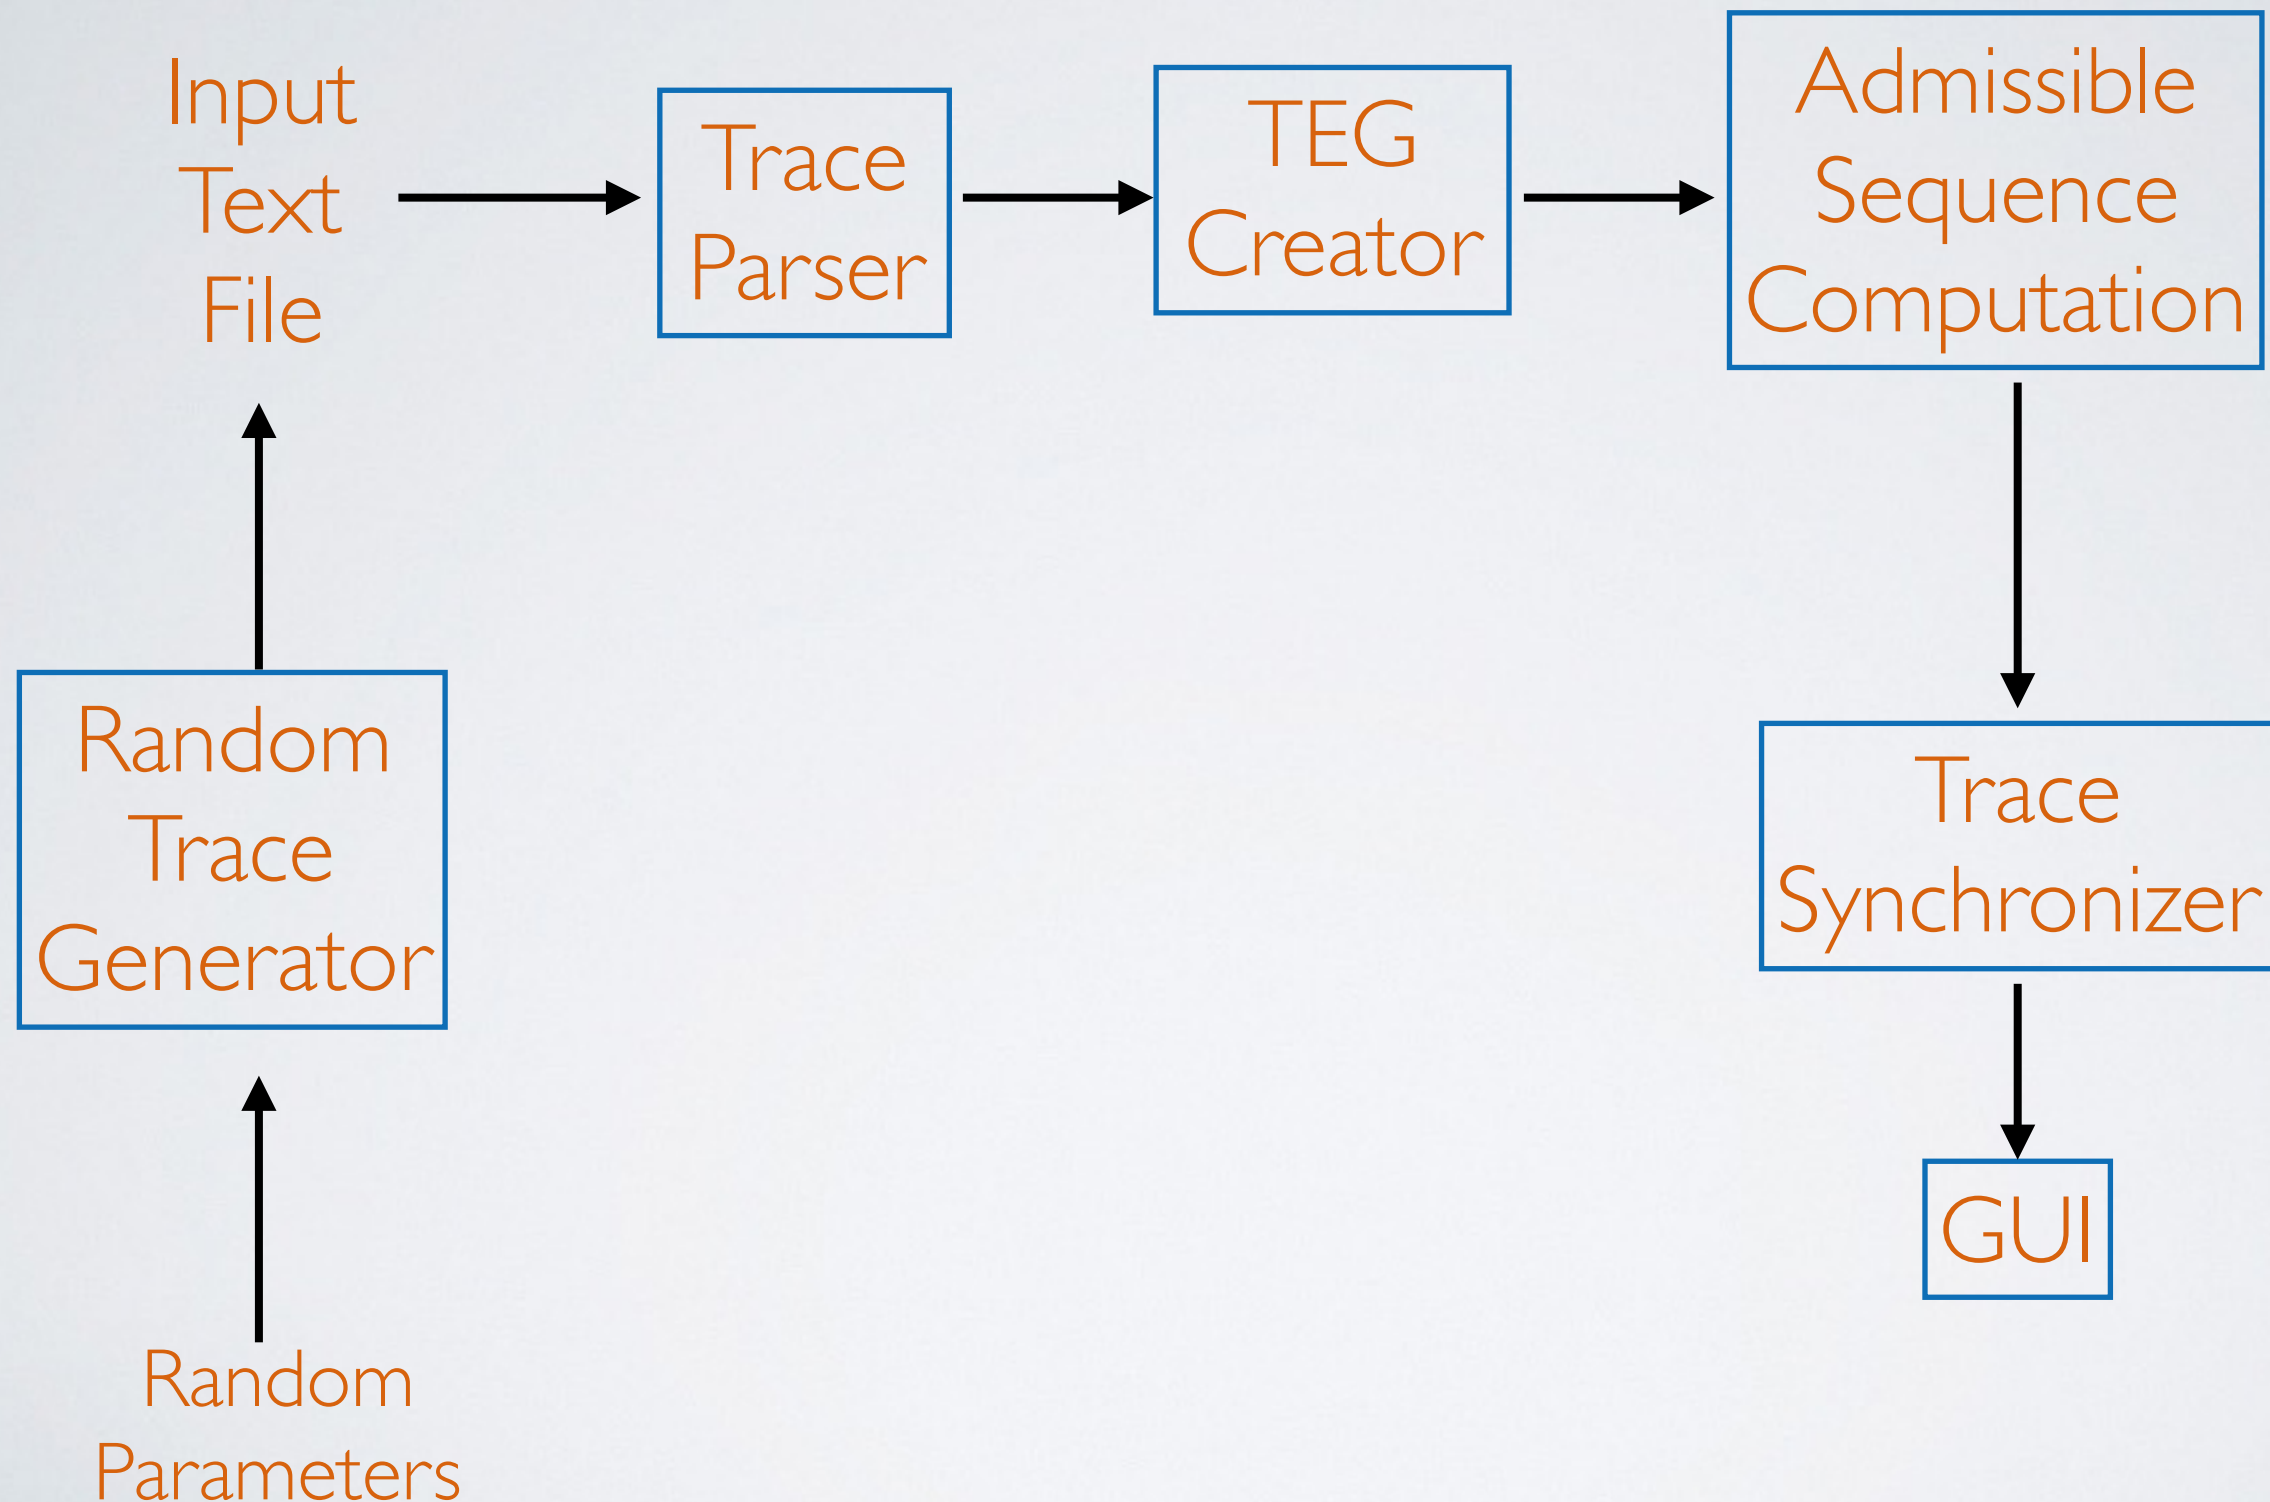


Admissible Sequence



Recompute Controlled Computation Trace

Java Framework



Admissible Sequence Algorithm

```
findAdmissibleSequence(trace):  
  graph = constructTEG(trace)  
  admissible_seq = NULL  
  min_length = INFINITY  
  for e in graph.initial_nodes  
    for f in graph.final_nodes  
      path = graph.shortest_path(e, f)  
      if (length(path) < min_length)  
        admissible_seq = path;  
        min_length = length(admissible_seq)  
      end if  
    end for  
  end for  
  return admissible_seq
```

Time Complexity: $O((nm)^3)$

Space Complexity: $O((nm)^2)$

Trace Synchronizer Algorithm

*addSynchrhonizationMessages (admissibleS_
seq) :*

*// add messages from set A in section
II*

for i in $[0, |\alpha|)$

for j in $(i, |\alpha|)$

if $(\alpha_i.\text{process} \neq \alpha_j.\text{process})$

 addSyncMessage(α_i , α_j)

end if

end for

end for

*// add messages from set B in section
II*

for i in $[0, |\alpha|-1)$


if $(\alpha_i.\text{process} \neq \alpha_{i+1}.\text{process})$

 addSyncMessage(α_{i+1} , $\alpha_i.\text{next}$)

end if


end for

True Event Graph




Original Computation

Admissible Sequence: $e1, f2, e4$

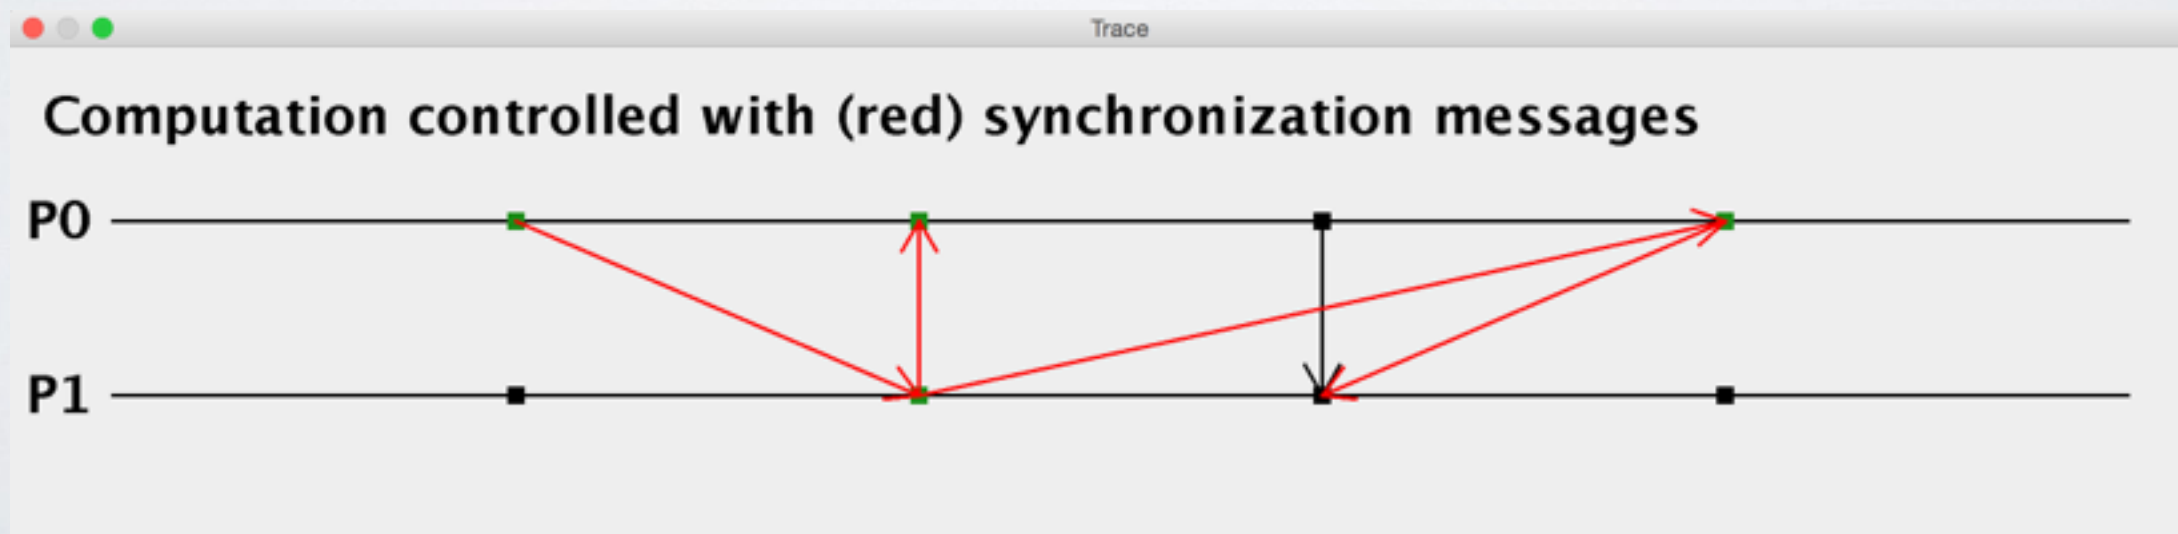
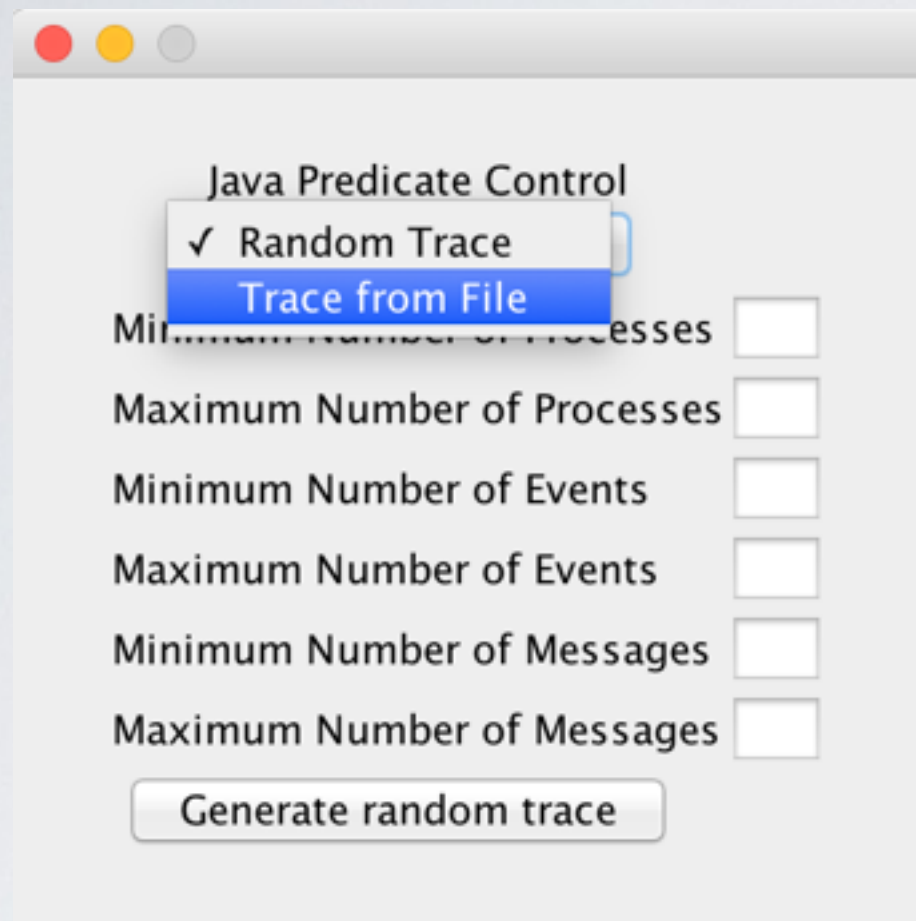


TEG



Controlled Computation

Java Framework



Java Framework

Java Predicate Control

Random Trace

Minimum Number of Processes

Maximum Number of Processes

Minimum Number of Events


Maximum Number of Events

Minimum Number of Messages

Maximum Number of Messages

Generate random trace

The predicate cannot be controllable in this computation



Java Framework

Java Predicate Control

Random Trace

Minimum Number of Processes

Maximum Number of Processes

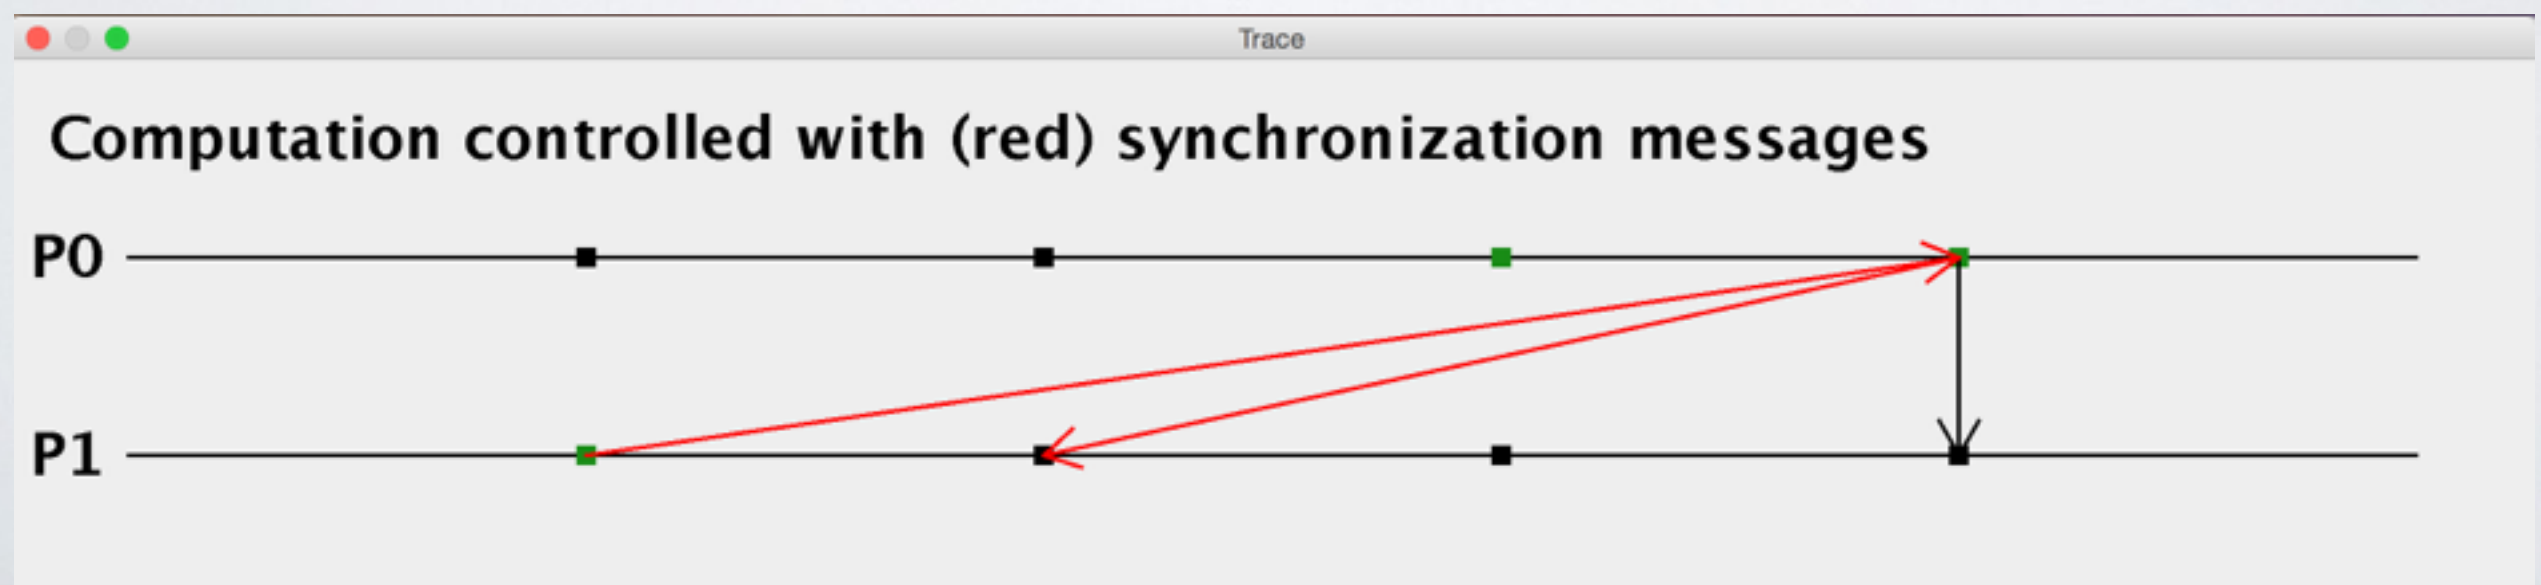
Minimum Number of Events

Maximum Number of Events

Minimum Number of Messages

Maximum Number of Messages

Generate random trace



Thank You!